

Speech Processing

Karen Livescu



Plan for this tutorial

- Yesterday: Intro to speech, historical tour of speech recognition research, speech recognition with hidden Markov models
- Today: Speech recognition with hybrid HMM/NNs and end-to-end recurrent neural networks, representation learning for speech
 - (Sorry, no language models)
- Lab exercise: Speech signals, speech recognition with HMMs and RNNs

Question from last time

Speech recognition error rates on low(er)-resource languages?

Question from last time

Speech recognition error rates on low(er)-resource languages?

Language	Language-specific	Joint	Joint + MTL
Bengali	19.1	16.8	16.5
Gujarati	26.0	18.0	18.2
Hindi	16.5	14.4	14.4
Kannada	35.4	34.5	34.6
Malayalam	44.0	36.9	36.7
Marathi	28.8	27.6	27.2
Tamil	13.3	10.7	10.6
Telugu	37.4	22.5	22.7
Urdu	29.5	26.8	26.7
Weighted Avg.	29.05	22.93	22.91

[Toshniwal+ 2018]

Model	Kazakh		Turkish		Haitian		Mongolian	
	WER	PER	WER	PER	WER	PER	WER	PER
Mono-lingual	55.9	40.9	53.1	36.2	49.0	36.9	58.2	45.2
Multi-lingual (MLing)	53.2	36.5	52.8	34.4	47.8	34.9	55.9	41.1
MLing & FineTuning (FT)	50.6	35.1	49.0	32.2	46.6	33.2	53.4	39.6
MLing + SWBD	52.3	36.6	51.3	33.0	45.8	33.9	54.5	40.2
MLing + SWBD & FT	48.2	33.5	48.7	31.9	44.3	31.9	51.5	37.8

[Dalmia+ 2018]

Outline

1 Speech recognition with hybrid HMM/NNs

2 Speech recognition with end-to-end neural models

3 Representation learning for speech

Hybrid HMM/NN models

Starting in the 1990s, long line of research by groups at ICSI Berkeley, IDIAP, and elsewhere (e.g., [Bourlard & Morgan 1994])

Hybrid HMM/NN models

Starting in the 1990s, long line of research by groups at ICSI Berkeley, IDIAP, and elsewhere (e.g., [Bourlard & Morgan 1994])

1. Train a frame-based neural network classifier of HMM states to produce a posterior $p(q|\mathbf{o})$, where q is the class and \mathbf{o} is a frame feature vector

Hybrid HMM/NN models

Starting in the 1990s, long line of research by groups at ICSI Berkeley, IDIAP, and elsewhere (e.g., [Bourlard & Morgan 1994])

1. Train a frame-based neural network classifier of HMM states to produce a posterior $p(q|\mathbf{o})$, where q is the class and \mathbf{o} is a frame feature vector
2. Convert $p(q|\mathbf{o})$ to something like an observation model (a “likelihood”): $p(\mathbf{o}|q) \propto \frac{p(q|\mathbf{o})}{p(q)}$

Hybrid HMM/NN models

Starting in the 1990s, long line of research by groups at ICSI Berkeley, IDIAP, and elsewhere (e.g., [Bourlard & Morgan 1994])

1. Train a frame-based neural network classifier of HMM states to produce a posterior $p(q|\mathbf{o})$, where q is the class and \mathbf{o} is a frame feature vector
2. Convert $p(q|\mathbf{o})$ to something like an observation model (a “likelihood”): $p(\mathbf{o}|q) \propto \frac{p(q|\mathbf{o})}{p(q)}$
3. Use the result in place of the observation model in an HMM

Hybrid models

Main idea is simple:

- Use a NN to produce a posterior for each class (= HMM state) given an input frame of acoustic features, $p(q|\mathbf{o})$
- Posterior is converted to a scaled likelihood via $p(\mathbf{o}|q) \propto \frac{p(q|\mathbf{o})}{p(q)}$

Hybrid models

Main idea is simple:

- Use a NN to produce a posterior for each class (= HMM state) given an input frame of acoustic features, $p(q|\mathbf{o})$
- Posterior is converted to a scaled likelihood via $p(\mathbf{o}|q) \propto \frac{p(q|\mathbf{o})}{p(q)}$

Where do the class labels for all frames come from?

Hybrid models

Main idea is simple:

- Use a NN to produce a posterior for each class (= HMM state) given an input frame of acoustic features, $p(q|\mathbf{o})$
- Posterior is converted to a scaled likelihood via $p(\mathbf{o}|q) \propto \frac{p(q|\mathbf{o})}{p(q)}$

Where do the class labels for all frames come from?

- Can be produced via a Viterbi alignment (“forced alignment”) using an existing HMM/GMM system

Hybrid models

Main idea is simple:

- Use a NN to produce a posterior for each class (= HMM state) given an input frame of acoustic features, $p(q|\mathbf{o})$
- Posterior is converted to a scaled likelihood via $p(\mathbf{o}|q) \propto \frac{p(q|\mathbf{o})}{p(q)}$

Where do the class labels for all frames come from?

- Can be produced via a Viterbi alignment (“forced alignment”) using an existing HMM/GMM system
- Or we can use “soft labels” = posteriors produced by running forward-backward using an existing HMM/GMM system
- (The latter makes sense if using cross-entropy loss)

How are HMMs used in speech research today?

- Hybrid HMM/NN models are close to (or at) state of the art for many benchmark data sets

How are HMMs used in speech research today?

- Hybrid HMM/NN models are close to (or at) state of the art for many benchmark data sets
- HMMs also used for **unsupervised** learning, e.g. discovering sound units in a low-resource language

How are HMMs used in speech research today?

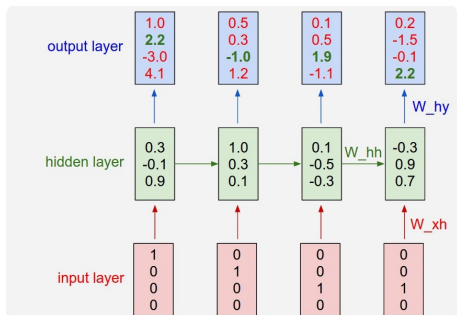
- Hybrid HMM/NN models are close to (or at) state of the art for many benchmark data sets
- HMMs also used for **unsupervised** learning, e.g. discovering sound units in a low-resource language
 - Given: a set of untranscribed speech
 - Train a single HMM on all of it
 - Look for repeated sequences of states to discover phones/words

Outline

- 1 Speech recognition with hybrid HMM/NNs
- 2 Speech recognition with end-to-end neural models
- 3 Representation learning for speech

Recurrent neural networks

Maintain a state vector in each frame, i.e. “remember” the past



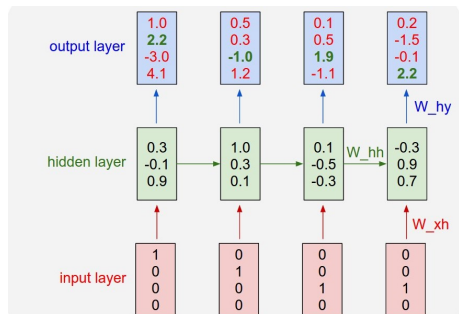
[Karpathy 2015]

$$\mathbf{h}_t = \sigma_h(\mathbf{W}_{xh}\mathbf{x}_t + \mathbf{W}_{hh}\mathbf{h}_{t-1} + \mathbf{b}_h)$$

$$\mathbf{y}_t = \sigma_y(\mathbf{W}_{hy}\mathbf{h}_t + \mathbf{b}_y)$$

Recurrent neural networks

Maintain a state vector in each frame, i.e. “remember” the past



[Karpathy 2015]

- Can be made deep, by stacking multiple layers of hidden states
- Can be made bidirectional, by combining a layer with forward connections and one with backward connections
- Variants: Long short-term memory (LSTM) networks, gated recurrent unit (GRU) networks

Using RNNs in hybrid HMM/NN models

If we add a softmax at the output layer and train with cross-entropy loss, the outputs approximate posterior probabilities of labels (HMM states):

$$y_t^i \approx p(q_t = i | \mathbf{x}_1, \dots, \mathbf{x}_t)$$

Using RNNs in hybrid HMM/NN models

If we add a softmax at the output layer and train with cross-entropy loss, the outputs approximate posterior probabilities of labels (HMM states):

$$y_t^i \approx p(q_t = i | \mathbf{x}_1, \dots, \mathbf{x}_t)$$

- Given this, we can produce scaled likelihoods and use them in a hybrid model, as before
- Or, we can drop the HMMs and use RNNs for **end-to-end** speech recognition...

End-to-end RNNs for speech recognition

Hybrid models involve a lot of machinery...

- Train a simple HMM/GMM recognizer
- Run Viterbi with HMM/GMM to get per-frame state labels
- Train neural network state classifier
- Compute state prior probabilities
- Scale classifier outputs by state priors to get scaled HMM observation model
- Train HMM/NN

End-to-end RNNs for speech recognition


Hybrid models involve a lot of machinery...

- Train a simple HMM/GMM recognizer
- Run Viterbi with HMM/GMM to get per-frame state labels
- Train neural network state classifier
- Compute state prior probabilities
- Scale classifier outputs by state priors to get scaled HMM observation model
- Train HMM/NN

Can we train an RNN to map directly from acoustic input sequence to output text sequence?

End-to-end RNNs for speech recognition


We could train an RNN to output a text label (word, character) at each frame...



[Karpathy 2015]

End-to-end RNNs for speech recognition

We could train an RNN to output a text label (word, character) at each frame...

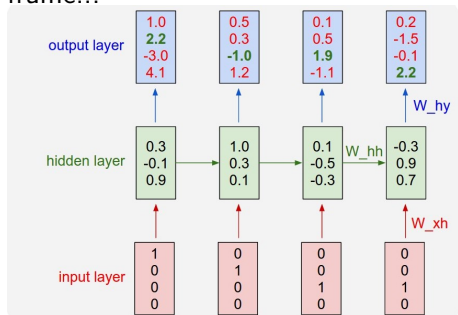


[Karpathy 2015]

- But what does this mean? Words and characters usually span many frames of speech

End-to-end RNNs for speech recognition

We could train an RNN to output a text label (word, character) at each frame...



[Karpathy 2015]


- But what does this mean? Words and characters usually span many frames of speech
- And the alignment between frames and characters/words is not simple

End-to-end RNNs for speech recognition

Two typical approaches:

- Encoder-decoder (“sequence-to-sequence”) models
- Connectionist temporal classification (CTC)


Encoder-decoder (“sequence-to-sequence”) RNNs



[Fig. credit smerity.com]

- Can be trained directly to optimize $p(\mathbf{y}|\mathbf{x})$ without aligning the input and output
- Input (acoustic frames) and output (characters/words) don't even have to operate at the same rate!
- Introduced for machine translation [Cho+ 2014, Sutskever+ 2014]

Encoder-decoder RNNs in more detail



[Fig. credit smerity.com]

“Vanilla” RNN encoder-decoder equations:

Encoder :

$$\mathbf{h}_t = \sigma_h(\mathbf{W}_{xh}\mathbf{x}_t + \mathbf{W}_{hh}\mathbf{h}_{t-1} + \mathbf{b}_h)$$

$$\mathbf{s} = \mathbf{h}_T = \mathbf{s}_0$$


Decoder :

$$\mathbf{s}_j = \sigma_s(\mathbf{W}_{ys}\mathbf{y}_{j-1} + \mathbf{W}_{ss}\mathbf{s}_{j-1} + \mathbf{b}_s)$$

$$\mathbf{f}_j = \text{softmax}(\mathbf{W}_{sy}\mathbf{s}_j + \mathbf{b}_y)$$

$$\hat{y}_j = \text{argmax } \mathbf{f}_j$$

Encoder-decoder RNNs in more detail



“Vanilla” RNN encoder-decoder equations:

Decoder :


$$\mathbf{s}_j = \sigma_s(\mathbf{W}_{ys}\mathbf{y}_{j-1} + \mathbf{W}_{ss}\mathbf{s}_{j-1} + \mathbf{b}_s)$$

$$\mathbf{f}_j = \text{softmax}(\mathbf{W}_{sy}\mathbf{s}_j + \mathbf{b}_y)$$

$$\hat{y}_j = \text{argmax } \mathbf{f}_j$$

- Here \mathbf{y}_j is a “one-hot” vector representing \hat{y}_j
- Interpretation: \mathbf{f}_{jd} is the probability of the next word being the d^{th} word in the vocabulary, given the previous words and the acoustic input

Encoder-decoder RNNs in more detail



“Vanilla” RNN encoder-decoder equations:

Decoder :


$$\mathbf{s}_j = \sigma_s(\mathbf{W}_{ys}\mathbf{y}_{j-1} + \mathbf{W}_{ss}\mathbf{s}_{j-1} + \mathbf{b}_s)$$

$$\mathbf{f}_j = \text{softmax}(\mathbf{W}_{sy}\mathbf{s}_j + \mathbf{b}_y)$$

$$\hat{y}_j = \text{argmax } \mathbf{f}_j$$

- Typical loss: cross-entropy (log loss)
 $-\log p(y_{1:J}|\mathbf{x}_{1:T}) = -\sum_{j=1}^J \log \mathbf{f}_{jd}$
- where $y_{1:J}$ = ground-truth label sequence corresponding to input sequence $\mathbf{x}_{1:T}$ and d is the index of y_j in the vocabulary

Encoder-decoder RNNs in more detail



Can be extended to a variety of types of encoder and decoder RNNs


- LSTM/GRU instead of vanilla RNN units
- Deep encoder, deep decoder (less typical)
- Bidirectional encoder

Attention models

- Basic encoder-decoder models represent the entire input sequence with a single vector
- That's a lot to ask of a single vector...
- Basic encoder-decoders don't work well out of the box. Much better when endowed with **attention**:


Attention models

- Basic encoder-decoder models represent the entire input sequence with a single vector
- That's a lot to ask of a single vector...
- Basic encoder-decoders don't work well out of the box. Much better when endowed with **attention**:



- In attention models, each decoder state depends on a weighted combination of encoder states (a “context vector”)
- These weights are an “attention vector”
- The attention vector is itself a function of the input and output, with learned parameters

Attention models



Example:

Decoder :

$$\mathbf{s}_j = \sigma_s(\mathbf{W}_{ys}\mathbf{y}_{j-1} + \mathbf{W}_{ss}\mathbf{s}_{j-1} + \mathbf{b}_s)$$

$$\mathbf{f}_j = \text{softmax}(\mathbf{W}_{sy}[\mathbf{c}_j; \mathbf{s}_j] + \mathbf{b}_y)$$

$$\hat{y}_j = \text{argmax } \mathbf{f}_j$$


Contextvector :

$$\mathbf{c}_j = \sum_{t=1}^T \alpha_{jt} \mathbf{h}_t$$

$$\alpha_j = \text{softmax}(\mathbf{u}_j)$$

$$\mathbf{u}_{jt} = \mathbf{h}_t^T \mathbf{s}_j$$

Attention models



Other ways of computing context vector:

$$\mathbf{c}_j = \sum_{t=1}^T \alpha_{jt} \mathbf{h}_t$$

$$\mathbf{alpha}_j = \text{softmax}(\mathbf{u}_j)$$

$$\mathbf{u}_{jt} = \mathbf{h}_t^T \mathbf{s}_j$$

$$\text{OR } \mathbf{u}_{jt} = \mathbf{v}^T \tanh(\mathbf{W}_1 \mathbf{h}_t + \mathbf{W}_2 \mathbf{s}_j + \mathbf{b}_a)$$

$$\text{OR } \mathbf{u}_{jt} = \mathbf{v}^T \tanh(\mathbf{W}_1 \mathbf{h}_t + \mathbf{W}_2 \mathbf{s}_j + \mathbf{W}_f \mathbf{f}_{jt} + \mathbf{b}_a)$$

$$\text{where } \mathbf{f}_j = \mathbf{F} * \mathbf{alpha}_{j-1}$$

The last is sometimes called “location-aware” or “convolutional” (attention)

Scheduled sampling [Bengio+ 2015]

During training, what should the input y_j to the decoder be?

Scheduled sampling [Bengio+ 2015]

During training, what should the input \mathbf{y}_j to the decoder be?

- To match test-time model, should set \mathbf{y}_j to a one-hot vector representing \hat{y}_{j-1}

Scheduled sampling [Bengio+ 2015]

During training, what should the input \mathbf{y}_j to the decoder be?

- To match test-time model, should set \mathbf{y}_j to a one-hot vector representing \hat{y}_{j-1}
- ... but then the training model would produce garbage output until it is trained well

Scheduled sampling [Bengio+ 2015]

During training, what should the input \mathbf{y}_j to the decoder be?

- To match test-time model, should set \mathbf{y}_j to a one-hot vector representing \hat{y}_{j-1}
- ... but then the training model would produce garbage output until it is trained well
- Or, we could set \mathbf{y}_j to a one-hot vector representing the ground-truth y_{j-1}

Scheduled sampling [Bengio+ 2015]

During training, what should the input \mathbf{y}_j to the decoder be?

- To match test-time model, should set \mathbf{y}_j to a one-hot vector representing \hat{y}_{j-1}
- ... but then the training model would produce garbage output until it is trained well
- Or, we could set \mathbf{y}_j to a one-hot vector representing the ground-truth y_{j-1}
- ... but then we are training and testing with different models


Scheduled sampling [Bengio+ 2015]

During training, what should the input \mathbf{y}_j to the decoder be?

- To match test-time model, should set \mathbf{y}_j to a one-hot vector representing \hat{y}_{j-1}
- ... but then the training model would produce garbage output until it is trained well
- Or, we could set \mathbf{y}_j to a one-hot vector representing the ground-truth y_{j-1}
- ... but then we are training and testing with different models
- Scheduled sampling: At each iteration of training, use the ground-truth label with some probability ϵ and the model's previous prediction with probability $1 - \epsilon$

Visualizing attention

FDHC0_SX209: Michael colored the bedroom wall with crayons.



[Chorowski+ 2015]

Attention models: Summary

Attention models have achieved state-of-the-art performance on some speech benchmarks. But:

- They are computationally demanding: Each output label considers entire input sequence
- But the alignment between the acoustics and labels is largely monotonic
- So maybe attention models are overkill?

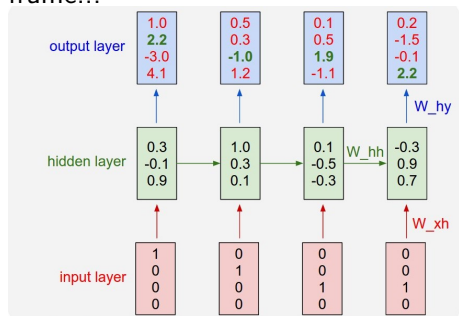
Rewind: End-to-end RNNs for speech recognition

Two typical approaches:

- Encoder-decoder (“sequence-to-sequence”) models
- Connectionist temporal classification (CTC)

Rewind: End-to-end RNNs for speech recognition


We could train an RNN to output a text label (word, character) at each frame...



[Karpathy 2015]

Rewind: End-to-end RNNs for speech recognition

We could train an RNN to output a text label (word, character) at each frame...




[Karpathy 2015]

- But what does this mean? Words and characters usually span many frames of speech

Rewind: End-to-end RNNs for speech recognition

We could train an RNN to output a text label (word, character) at each frame...



[Karpathy 2015]

- But what does this mean? Words and characters usually span many frames of speech
- And the alignment between frames and characters/words is not simple, and often ambiguous. Consider the word **through**: Which frames does the **g** correspond to?

But given the mostly monotonic alignment between frames and labels, maybe we were too dismissive?

Connectionist temporal classification (CTC) [Graves+

2006]


CTC modifies the per-frame RNN labeler idea in two key ways:

- An extra “blank” label ϵ
- A mapping from frame-level label sequences to true label sequences

Connectionist temporal classification (CTC) [Graves+

2006]

From <https://distill.pub/2017/ctc/> [Hannun 2017]:



h	h	h	h	h	h	h	h	h	h
e	e	e	e	e	e	e	e	e	e
l	l	l	l	l	l	l	l	l	l
o	o	o	o	o	o	o	o	o	o
€	€	€	€	€	€	€	€	€	€

RNN with softmax output layer produces a posterior probability for each label + ϵ

Basic (“greedy”) CTC decoding

- RNN with softmax output layer produces a posterior probability for each label $+ \epsilon$
- At each time frame, output the most likely frame label
- Finally, map frame labels to “collapsed” label sequence as follows:

h h e ϵ ϵ l l l ϵ l l o

First, merge repeat characters.

h e ϵ l ϵ l o

Then, remove any ϵ tokens.

h e l l o

The remaining characters are the output.

h e l l o

[Hannun 2017]

CTC training

Given a sequence X of N acoustic frames and a corresponding label sequence Y with $L < N$ labels, consider the set of all of the valid frame label sequences (“alignments”) $\mathcal{A}_{X,Y}$

Valid Alignments

€ c c € a t

c c a a t t

c a € € € t

Invalid Alignments

c € c € a t

c c a a t

c € € € | t t

corresponds to
 $Y = [c, c, a, t]$

has length 5

missing the 'a'

[Hannun 2017]

(Example for the word **cat**)

CTC training

Given a sequence X of T acoustic frames and a corresponding label sequence Y with $L < N$ labels, e.g. the word **cat**, consider the set of all of the valid frame label sequences (“alignments”) $\mathcal{A}_{X,Y}$.

Then the CTC loss is a *marginal log loss*:

$$-\log p(Y|X) = \sum_{A \in \mathcal{A}_{X,Y}} \prod_{t=1}^T p_t(a_t|X)$$

where $p_t(a_t|X)$ is the softmax output of the RNN at frame t

CTC training


CTC loss:

$$-\log p(Y|X) = \sum_{A \in \mathcal{A}_{X,Y}} \prod_{t=1}^T p_t(a_t|X)$$

Looks hard to backprop, but it turns out to be equivalent to a forward-backward-like HMM algorithm!

CTC posterior visualization

CTC posteriors vs. posteriors from an RNN trained with frame-level log loss (e.g. for a hybrid HMM/NN):



[Graves+ 2006]

CTC not so different from HMMs...

Putting HMM into CTC-like notation (A = state sequence):

$$p(X|Y) = \sum_{A \in \mathcal{A}} p(X, A|Y)$$

Dropping the conditioning on Y :

$$p(X) = \sum_{A \in \mathcal{A}} \prod_{t=1}^T p(\mathbf{x}_t|a_t)p(a_t|a_{t-1})$$

Suppose transition probabilities are uniform:

$$p(X) \propto \sum_{A \in \mathcal{A}} \prod_{t=1}^T p(\mathbf{x}_t|a_t)$$

CTC not so different from HMMs...

$$p(X) \propto \sum_{A \in \mathcal{A}} \prod_{t=1}^T p(\mathbf{x}_t | a_t)$$

2 differences from CTC:

- $p(\mathbf{x}_t | a_t)$ vs. $p(a_t | \mathbf{x}_t)$
- Definition of \mathcal{A}

Rewrite using Bayes rule (as we did for hybrid models):

$$p(X) \propto \sum_{A \in \mathcal{A}} \prod_{t=1}^T p(a_t | \mathbf{x}_t) / p(a_t)$$


Assuming uniform priors of the labels (states):

$$p(X) \propto \sum_{A \in \mathcal{A}} \prod_{t=1}^T p(a_t | \mathbf{x}_t)$$

So computing the marginal probability in CTC is just like computing likelihood in HMMs, hence forward-backward algorithm!

Equivalent HMM state diagram for CTC

Assuming the ground-truth sequence "a b":



Connection between CTC and encoder-decoder models

- Encoder = all but last layer of the RNN
- Decoder = softmax + label collapsing function

CTC: Summary

- CTC-based models are state-of-the-art on many speech benchmarks
- Much faster than attention models
- Tend to require more data to train well

End-to-end neural speech recognition: Addendum

This is a very active area of research

- Models we haven't discussed: transformers, dilated convolution-based models, ...
- Issues we haven't discussed:
 - Choice of output labels (words, characters, other sub-word units)
 - How to combine with a language model

Outline

- 1 Speech recognition with hybrid HMM/NNs
- 2 Speech recognition with end-to-end neural models
- 3 Representation learning for speech

“Tandem” models

- Grew out of early work on hybrid HMM/NN models
- Developed at ICSI Berkeley (e.g., Hermansky et al. 2000)
- Main idea: Never mind the whole posterior conversion business; use the NN classifier outputs as features!

“Tandem” models

- Grew out of early work on hybrid HMM/NN models
- Developed at ICSI Berkeley (e.g., Hermansky et al. 2000)
- Main idea: Never mind the whole posterior conversion business; use the NN classifier outputs as features!
 - Then use standard HMM/GMMs with these features as inputs

“Tandem” models


- Grew out of early work on hybrid HMM/NN models
- Developed at ICSI Berkeley (e.g., Hermansky et al. 2000)
- Main idea: Never mind the whole posterior conversion business; use the NN classifier outputs as features!
 - Then use standard HMM/GMMs with these features as inputs
 - $\mathbf{o}' = [y_1(\mathbf{o}) \ y_2(\mathbf{o}) \ \dots y_n(\mathbf{o})]$

“Tandem” models

- Grew out of early work on hybrid HMM/NN models
- Developed at ICSI Berkeley (e.g., Hermansky et al. 2000)
- Main idea: Never mind the whole posterior conversion business; use the NN classifier outputs as features!
 - Then use standard HMM/GMMs with these features as inputs
 - $\mathbf{o}' = [y_1(\mathbf{o}) \ y_2(\mathbf{o}) \ \dots \ y_n(\mathbf{o})]$
 - If the $y_i(\mathbf{o})$ represent probabilities, then we typically take their logs:
 $\mathbf{o}' = [\log(f_1(\mathbf{o})) \ \log(f_2(\mathbf{o})) \ \dots]$

Tandem models

Alternatively, use outputs from a lower layer, and make that layer narrow (a “bottleneck layer”) to reduce dimensionality



[Gehring+ 2013]

Tandem models: More details

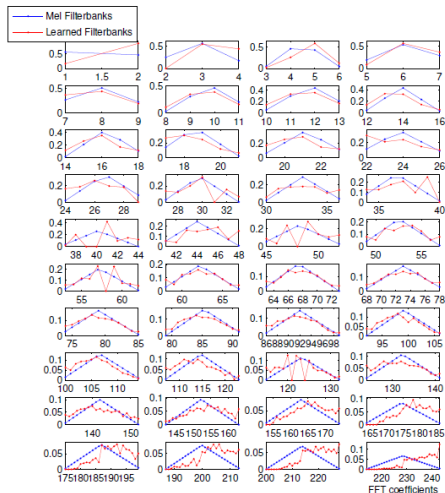
- These features are often appended to the original features, so the new feature vector is $[\mathbf{o} \ \mathbf{o}']$ (hence, “tandem”!)

Tandem models: More details

- These features are often appended to the original features, so the new feature vector is $[\mathbf{o} \ \mathbf{o}']$ (hence, “tandem”!)
- Typically, the input is a concatenation of acoustic vectors over a window of 7-20 frames around the current frame (very high-dimensional!)

Tandem models: Learned representations

When using spectrograms as input representation, the learned filters are similar to the triangular filters of the inner ear:



Tandem models: Summary

- Tandem models were state-of-the-art around the late 2000s
- Both supervised and unsupervised “bottleneck” representations were attempted, but only the supervised ones performed well
- Extensions to multilingual learning of bottleneck features have been quite successful

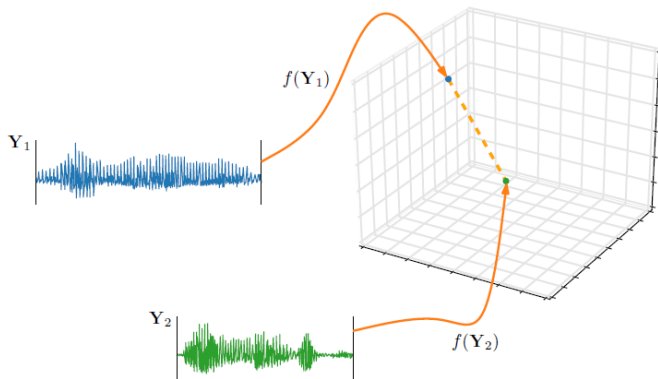
Current research: Acoustic word embeddings

Frame representations are useful, but for some tasks we would like a representation of whole words (or other units) of arbitrary duration

Current research: Acoustic word embeddings

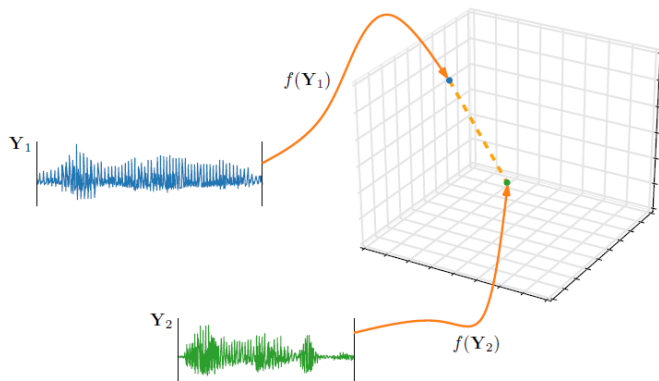
Frame representations are useful, but for some tasks we would like a representation of whole words (or other units) of arbitrary duration

- Acoustic word embeddings: map from a spoken word to a vector
- “Spoken word” = speech signal of arbitrary length corresponding to a word



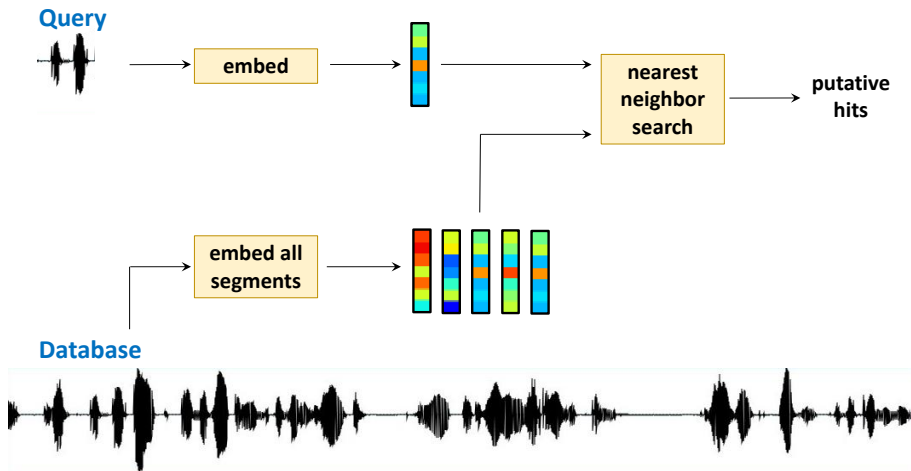
Applications of acoustic word embeddings

- Speech search [Parada+ 2015, Settle+ 2017, Audhkhasi+ 2017]
- Whole-word speech recognition [Maas+ 2012, Bengio & Heigold 2014, Settle+ 2019]
- Spoken term discovery [Kamper+ 2014-2018]



[Figure credit: Herman Kamper]


Query-by-example search with acoustic word embeddings [Settle+ 2017]



[Figure credit: Herman Kamper]

Neural embeddings: RNN-based [SLT 2016]

- **Input:** MFCCs (without padding)
- **Model:** n_{rec} recurrent + n_{full} fully connected layers
- **Embedding** is activation vector of final fully connected layer



Training objectives

Word classifier log loss

- Add a softmax layer to predict word w
- $l(\mathbf{x}, w) = \log p(w|\mathbf{x})$

Training objectives

Word classifier log loss

- Add a softmax layer to predict word w
- $l(\mathbf{x}, w) = \log p(w|\mathbf{x})$

Contrastive (triplet) loss

- Bring together same-word pairs, separate different ones

$$l(\mathbf{x}_1, \mathbf{x}_2) = \max\{0, m + d_{\cos}(\mathbf{x}_1, \mathbf{x}_2) - d_{\cos}(\mathbf{x}_1, \mathbf{x}^-)\}$$

where \mathbf{x}^- = random (or hard) negative example, m = margin


- Weaker supervision (no word labels, only same-word pairs)

Visualization: RNN embeddings

2-dimensional t-SNE embeddings [van der Maaten & Hinton 2008]

\triangle = word types seen at training time

\circ = not seen at training time



These embeddings far outperform a standard (dynamic time warping-based) approach to query-by-example search

Joint learning of acoustic + written word embeddings

[He+ 2017, Settle+ 2019, Collobert+ 2019]

Motivation:

- Learn better acoustic embeddings by relating them to a written character sequence
- Some tasks involve “distances” between speech segments and written words
 - Spoken term detection (“Query-by-text”)
 - Automatic speech recognition

“Barack Obama”

?

=



Joint learning of acoustic + written word embeddings [He+ 2017, Settle+ 2019, Collobert+ 2019]

Motivation:

- Learn better acoustic embeddings by relating them to a written character sequence
- Some tasks involve “distances” between speech segments and written words
 - Spoken term detection (“Query-by-text”)
 - Automatic speech recognition

Approach: Learn a pair of RNN-based embedding functions

- Acoustic word embedding (speech \rightarrow vector)
- Acoustically grounded word embedding (character sequence \rightarrow vector)

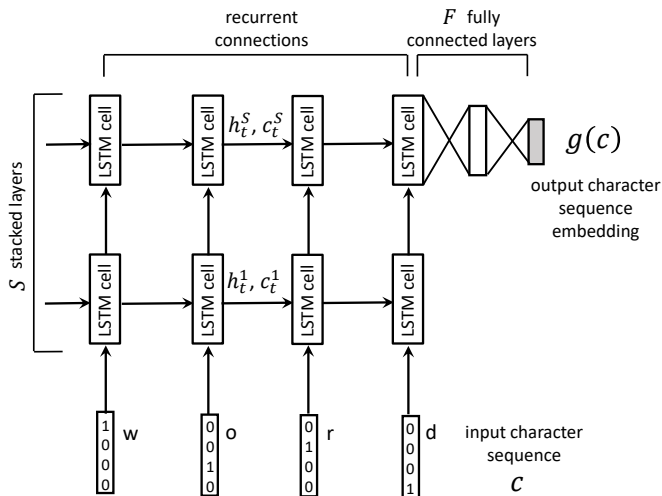
“Barack Obama”

?

=



Character RNN-based acoustically grounded word embedding



Joint learning of acoustic and acoustically grounded word embeddings

Given a matched (acoustic, written) word pair (\mathbf{x}, \mathbf{c})

$$l_0(\mathbf{x}, \mathbf{c}) = \max\{0, m + d_{\cos}(\mathbf{x}, \mathbf{c}) - d_{\cos}(\mathbf{x}, \mathbf{c}^-)\}$$

Joint learning of acoustic and acoustically grounded word embeddings

Given a matched (acoustic, written) word pair (\mathbf{x}, \mathbf{c})

$$l_0(\mathbf{x}, \mathbf{c}) = \max\{0, m + d_{\cos}(\mathbf{x}, \mathbf{c}) - d_{\cos}(\mathbf{x}, \mathbf{c}^-)\}$$

$$l_1(\mathbf{x}, \mathbf{c}) = \max\{0, m + d_{\cos}(\mathbf{x}, \mathbf{c}) - d_{\cos}(\mathbf{c}^-, \mathbf{c})\}$$

Joint learning of acoustic and acoustically grounded word embeddings

Given a matched (acoustic, written) word pair (\mathbf{x}, \mathbf{c})

$$l_0(\mathbf{x}, \mathbf{c}) = \max\{0, m + d_{\cos}(\mathbf{x}, \mathbf{c}) - d_{\cos}(\mathbf{x}, \mathbf{c}^-)\}$$

$$l_1(\mathbf{x}, \mathbf{c}) = \max\{0, m + d_{\cos}(\mathbf{x}, \mathbf{c}) - d_{\cos}(\mathbf{c}^-, \mathbf{c})\}$$

$$l_2(\mathbf{x}, \mathbf{c}) = \max\{0, m + d_{\cos}(\mathbf{x}, \mathbf{c}) - d_{\cos}(\mathbf{x}^-, \mathbf{c})\}$$

$$l_3(\mathbf{x}, \mathbf{c}) = \max\{0, m + d_{\cos}(\mathbf{x}, \mathbf{c}) - d_{\cos}(\mathbf{x}, \mathbf{x}^-)\}$$

Joint learning of acoustic and acoustically grounded word embeddings

Given a matched (acoustic, written) word pair (\mathbf{x}, \mathbf{c})

$$l_0(\mathbf{x}, \mathbf{c}) = \max\{0, m + d_{\cos}(\mathbf{x}, \mathbf{c}) - d_{\cos}(\mathbf{x}, \mathbf{c}^-)\}$$

$$l_1(\mathbf{x}, \mathbf{c}) = \max\{0, m + d_{\cos}(\mathbf{x}, \mathbf{c}) - d_{\cos}(\mathbf{c}^-, \mathbf{c})\}$$


$$l_2(\mathbf{x}, \mathbf{c}) = \max\{0, m + d_{\cos}(\mathbf{x}, \mathbf{c}) - d_{\cos}(\mathbf{x}^-, \mathbf{c})\}$$

$$l_3(\mathbf{x}, \mathbf{c}) = \max\{0, m + d_{\cos}(\mathbf{x}, \mathbf{c}) - d_{\cos}(\mathbf{x}, \mathbf{x}^-)\}$$

Variants:

- Weighted combination of these losses
- Cost-sensitive margin that scales with orthographic distance

Visualization of acoustic + acoustically grounded word embeddings



Whole-word speech recognition with joint acoustic/written word embeddings

Whole-word ASR [Audhkhasi+ 2017]

- Output labels are whole words
- Final layer weights represent a word embedding matrix
- Many rare words \implies many rows are learned very poorly
- **Idea:** Pre-train the RNN with acoustic word embeddings and the softmax layer weights with jointly trained written word embeddings
- This improves speech recognition performance [Settle+ 2019, Collobert+ 2019]
- **Bonus:** Can recognize previously unseen words

Other exciting current speech research...

- Unsupervised representation learning (is starting to work!)
- Learning speech representations from multimodal data
- Learning **semantic** speech representations from visually grounded data
- Multilingual models
- End-to-end spoken language understanding
- Speech generation

THE END